

<https://towardsdatascience.com/machine-learning-basics-logistic-regression-890ef5e3a272> (<https://towardsdatascience.com/machine-learning-basics-logistic-regression-890ef5e3a272>)

## Importing the Libraries

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

## Importing the dataset

```
In [2]: dataset = pd.read_csv('https://raw.githubusercontent.com/mk-gurucharan
X = dataset.iloc[:, [0, 1]].values
y = dataset.iloc[:, 2].values
```

```
In [3]: dataset.head(5)
```

Out[3]:

	DMV_Test_1	DMV_Test_2	Results
0	34.623660	78.024693	0
1	30.286711	43.894998	0
2	35.847409	72.902198	0
3	60.182599	86.308552	1
4	79.032736	75.344376	1

## Splitting the dataset into the Training set and Test set

In this the `test_size=0.25` denotes that 25% of the data will be kept as the Test set and the remaining 75% will be used for training as the Training set.

```
In [4]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size =
```

## Feature Scaling

This is an additional step that is used to normalize the data within a particular range. It also aids in speeding up the calculations. As the data is widely varying, we use this function to limit the range of the data within a small limit ( -2,2). For example, the score 62.0730638 is normalized to -0.21231162 and the score 96.51142588 is normalized to 1.55187648. In this way, the scores of X\_train and X\_test are normalized to a smaller range.

```
In [6]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

## Training the Logistic Regression model on the Training Set

In this step, the class LogisticRegression is imported and is assigned to the variable "classifier". The classifier.fit() function is fitted with X\_train and Y\_train on which the model will be trained.

```
In [7]: from sklearn.linear_model import LogisticRegression
classifier = LogisticRegression()
classifier.fit(X_train, y_train)
```

```
Out[7]: LogisticRegression()
```

## Predicting the Test set results

In this step, the classifier.predict() function is used to predict the values for the Test set and the values are stored to the variable y\_pred.

```
In [8]: y_pred = classifier.predict(X_test)
y_pred
```

```
Out[8]: array([1, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0,
        , 0,
        , 1, 1, 0])
```

## Confusion Matrix and Accuracy

The confusion matrix is a table that is used to show the number of correct and incorrect predictions on a classification problem when the real values of the Test Set are known. It is of the format

The True values are the number of correct predictions made.

```
In [9]: from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
from sklearn.metrics import accuracy_score
print ("Accuracy : ", accuracy_score(y_test, y_pred))
cm
```

Accuracy : 0.88

```
Out[9]: array([[11,  0],
               [ 3, 11]])
```

## Comparing the Real Values with Predicted Values

*In this step, a Pandas DataFrame is created to compare the classified values of both the original Test set (y\_test) and the predicted results (y\_pred).*

```
In [10]: df = pd.DataFrame({'Real Values':y_test, 'Predicted Values':y_pred})
df
```

Out[10]:

	Real Values	Predicted Values
0	1	1
1	0	0
2	0	0
3	0	0
4	1	1
5	1	1
6	1	0
7	1	1
8	0	0
9	1	1
10	0	0
11	0	0
12	0	0
13	1	1
14	1	0
15	1	1
16	0	0
17	1	1
18	1	0
19	1	1
20	0	0
21	0	0
22	1	1
23	1	1
24	0	0


## Visualising the Results

In this last step, we visualize the results of the Logistic Regression model on a graph that is plotted along with the two regions

```
In [11]: from matplotlib.colors import ListedColormap
X_set, y_set = X_test, y_test
X1, X2 = np.meshgrid(np.arange(start = X_set[:, 0].min() - 1, stop = X
                        np.arange(start = X_set[:, 1].min() - 1, stop = X
plt.contourf(X1, X2, classifier.predict(np.array([X1.ravel(), X2.ravel
alpha = 0.75, cmap = ListedColormap(('red', 'green'))))
plt.xlim(X1.min(), X1.max())
plt.ylim(X2.min(), X2.max())
for i, j in enumerate(np.unique(y_set)):
    plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1],
                c = ListedColormap(('red', 'green'))(i), label = j)
plt.title('Logistic Regression')
plt.xlabel('DMV_Test_1')
plt.ylabel('DMV_Test_2')
plt.legend()
plt.show()
```

\*c\* argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence in case its length matches with \*x\* & \*y\*. Please use the \*color\* keyword-argument or provide a 2-D array with a single row if you intend to specify the same RGB or RGBA value for all points.

\*c\* argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence in case its length matches with \*x\* & \*y\*. Please use the \*color\* keyword-argument or provide a 2-D array with a single row if you intend to specify the same RGB or RGBA value for all points.



In [ ]:

